

object-detection-yolo

April 2, 2020

Object detection and classification with YOLOv2

Project created during the Deep Learning Specialization course on www.coursera.org

```
[1]: import warnings
warnings.filterwarnings(action='ignore')
import os
import matplotlib.pyplot as plt
from matplotlib.pyplot import imshow
import numpy as np
import imageio
import tensorflow as tf
from PIL import Image, ImageDraw, ImageFont
from keras import backend as K
import colorsys
import random
from io import BytesIO
from keras.models import load_model

%matplotlib inline

Using TensorFlow backend.
```

Input image preprocessing

```
[2]: # Preprocess input images to match model input size

def preprocess_image(img_path, model_image_size):

    image = Image.open(img_path)

    # Retrieve image shape (necessary to scale predicted bounding boxes later on)
    image_shape = image.size[::-1]

    # Resize image to correspond to model image size with BICUBIC interpolation
    ↪mode (cubic spline interpolation)
    resized_image = image.resize(model_image_size, Image.BICUBIC)
```

```

# Convert image to a numpy array
image_data = np.array(resized_image, dtype='float32')

# Normalize image
image_data /= 255.

# Add batch dimension
image_data = np.expand_dims(image_data, 0)

# Display the image
plt.show()

return image, image_data, image_shape

```

Defining anchor boxes

```
[3]: # Load anchor boxes dimensions (predefined)
```

```

def read_anchors(anchors_path):
    with open(anchors_path) as f:
        anchors = f.readline()
        anchors = [float(x) for x in anchors.split(',')]
        anchors = np.array(anchors).reshape(-1, 2)
    return anchors

```

```
[4]: anchors = read_anchors('./yad2k/model_data/yolo_anchors.txt')
anchors
```

```
[4]: array([[0.57273 , 0.677385],
            [1.87446 , 2.06253 ],
            [3.33843 , 5.47434 ],
            [7.88282 , 3.52778 ],
            [9.77052 , 9.16828 ]])
```

Defining labels

```
[5]: # Load labels from COCO dataset
```

```

def read_classes(classes_path):
    with open(classes_path) as f:
        class_names = f.readlines()
        class_names = [c.strip() for c in class_names]
    return class_names

```

```
[6]: class_names = read_classes('./yad2k/model_data/coco_classes.txt')
    class_names[:10]

[6]: ['person',
      'bicycle',
      'car',
      'motorbike',
      'aeroplane',
      'bus',
      'train',
      'truck',
      'boat',
      'traffic light']

[7]: # Generate colors to represent classes on the image (for drawing bounding boxes)

def generate_colors(class_names):

    hsv_tuples = [(x / len(class_names), 1., 1.) for x in
↪range(len(class_names))]
    colors = list(map(lambda x: colorsys.hsv_to_rgb(*x), hsv_tuples))
    colors = list(map(lambda x: (int(x[0] * 255), int(x[1] * 255), int(x[2] *
↪255)), colors))

    # Fixed seed for consistent colors across runs
    random.seed(10101)

    # Shuffle colors to decorrelate adjacent classes
    random.shuffle(colors)

    # Reset seed to default
    random.seed(None)

    return colors
```

Loading pretrained YOLOv2

```
[8]: sess = K.get_session()

[9]: # Load pre-trained model from the official YOLO site, according to instructions
↪in this repo:
    # https://github.com/allanzelener/YAD2K
    model = load_model('./yad2k/model_data/yolo.h5')

WARNING:tensorflow:From /usr/local/lib/python3.6/site-
packages/keras/backend/tensorflow_backend.py:4070: The name tf.nn.max_pool is
deprecated. Please use tf.nn.max_pool2d instead.
```

WARNING:tensorflow:From /Users/alicjamazur/Desktop/dl/keras/utils_yolo/yad2k/yad2k/models/keras_yolo.py:32: The name tf.space_to_depth is deprecated. Please use tf.compat.v1.space_to_depth instead.

[10]: # Show model summary

model.summary()

Model: "model_1"

| Layer (type) | Output Shape | Param # | Connected to |
|--|-----------------------|---------|-----------------------------|
| input_1 (InputLayer) | (None, 608, 608, 3) | 0 | |
| conv2d_1 (Conv2D) | (None, 608, 608, 32) | 864 | input_1[0][0] |
| batch_normalization_1 (BatchNormalizatio | (None, 608, 608, 32) | 128 | conv2d_1[0][0] |
| leaky_re_lu_1 (LeakyReLU) | (None, 608, 608, 32) | 0 | batch_normalization_1[0][0] |
| max_pooling2d_1 (MaxPooling2D) | (None, 304, 304, 32) | 0 | leaky_re_lu_1[0][0] |
| conv2d_2 (Conv2D) | (None, 304, 304, 64) | 18432 | max_pooling2d_1[0][0] |
| batch_normalization_2 (BatchNormalizatio | (None, 304, 304, 64) | 256 | conv2d_2[0][0] |
| leaky_re_lu_2 (LeakyReLU) | (None, 304, 304, 64) | 0 | batch_normalization_2[0][0] |
| max_pooling2d_2 (MaxPooling2D) | (None, 152, 152, 64) | 0 | leaky_re_lu_2[0][0] |
| conv2d_3 (Conv2D) | (None, 152, 152, 128) | 73728 | |

```

max_pooling2d_2[0][0]
-----
batch_normalization_3 (BatchNor (None, 152, 152, 128 512      conv2d_3[0][0]
-----
leaky_re_lu_3 (LeakyReLU)      (None, 152, 152, 128 0
batch_normalization_3[0][0]
-----
conv2d_4 (Conv2D)              (None, 152, 152, 64) 8192
leaky_re_lu_3[0][0]
-----
batch_normalization_4 (BatchNor (None, 152, 152, 64) 256      conv2d_4[0][0]
-----
leaky_re_lu_4 (LeakyReLU)      (None, 152, 152, 64) 0
batch_normalization_4[0][0]
-----
conv2d_5 (Conv2D)              (None, 152, 152, 128 73728
leaky_re_lu_4[0][0]
-----
batch_normalization_5 (BatchNor (None, 152, 152, 128 512      conv2d_5[0][0]
-----
leaky_re_lu_5 (LeakyReLU)      (None, 152, 152, 128 0
batch_normalization_5[0][0]
-----
max_pooling2d_3 (MaxPooling2D) (None, 76, 76, 128) 0
leaky_re_lu_5[0][0]
-----
conv2d_6 (Conv2D)              (None, 76, 76, 256) 294912
max_pooling2d_3[0][0]
-----
batch_normalization_6 (BatchNor (None, 76, 76, 256) 1024      conv2d_6[0][0]
-----
leaky_re_lu_6 (LeakyReLU)      (None, 76, 76, 256) 0
batch_normalization_6[0][0]
-----
conv2d_7 (Conv2D)              (None, 76, 76, 128) 32768

```

```

leaky_re_lu_6[0][0]
-----
batch_normalization_7 (BatchNor (None, 76, 76, 128) 512 conv2d_7[0][0]
-----
leaky_re_lu_7 (LeakyReLU) (None, 76, 76, 128) 0
batch_normalization_7[0][0]
-----
conv2d_8 (Conv2D) (None, 76, 76, 256) 294912
leaky_re_lu_7[0][0]
-----
batch_normalization_8 (BatchNor (None, 76, 76, 256) 1024 conv2d_8[0][0]
-----
leaky_re_lu_8 (LeakyReLU) (None, 76, 76, 256) 0
batch_normalization_8[0][0]
-----
max_pooling2d_4 (MaxPooling2D) (None, 38, 38, 256) 0
leaky_re_lu_8[0][0]
-----
conv2d_9 (Conv2D) (None, 38, 38, 512) 1179648
max_pooling2d_4[0][0]
-----
batch_normalization_9 (BatchNor (None, 38, 38, 512) 2048 conv2d_9[0][0]
-----
leaky_re_lu_9 (LeakyReLU) (None, 38, 38, 512) 0
batch_normalization_9[0][0]
-----
conv2d_10 (Conv2D) (None, 38, 38, 256) 131072
leaky_re_lu_9[0][0]
-----
batch_normalization_10 (BatchNo (None, 38, 38, 256) 1024 conv2d_10[0][0]
-----
leaky_re_lu_10 (LeakyReLU) (None, 38, 38, 256) 0
batch_normalization_10[0][0]
-----
conv2d_11 (Conv2D) (None, 38, 38, 512) 1179648

```

```

leaky_re_lu_10[0][0]
-----
batch_normalization_11 (BatchNo (None, 38, 38, 512) 2048 conv2d_11[0][0]
-----
leaky_re_lu_11 (LeakyReLU) (None, 38, 38, 512) 0
batch_normalization_11[0][0]
-----
conv2d_12 (Conv2D) (None, 38, 38, 256) 131072
leaky_re_lu_11[0][0]
-----
batch_normalization_12 (BatchNo (None, 38, 38, 256) 1024 conv2d_12[0][0]
-----
leaky_re_lu_12 (LeakyReLU) (None, 38, 38, 256) 0
batch_normalization_12[0][0]
-----
conv2d_13 (Conv2D) (None, 38, 38, 512) 1179648
leaky_re_lu_12[0][0]
-----
batch_normalization_13 (BatchNo (None, 38, 38, 512) 2048 conv2d_13[0][0]
-----
leaky_re_lu_13 (LeakyReLU) (None, 38, 38, 512) 0
batch_normalization_13[0][0]
-----
max_pooling2d_5 (MaxPooling2D) (None, 19, 19, 512) 0
leaky_re_lu_13[0][0]
-----
conv2d_14 (Conv2D) (None, 19, 19, 1024) 4718592
max_pooling2d_5[0][0]
-----
batch_normalization_14 (BatchNo (None, 19, 19, 1024) 4096 conv2d_14[0][0]
-----
leaky_re_lu_14 (LeakyReLU) (None, 19, 19, 1024) 0
batch_normalization_14[0][0]
-----
conv2d_15 (Conv2D) (None, 19, 19, 512) 524288

```

```

leaky_re_lu_14[0][0]
-----
batch_normalization_15 (BatchNo (None, 19, 19, 512) 2048 conv2d_15[0][0]
-----
leaky_re_lu_15 (LeakyReLU) (None, 19, 19, 512) 0
batch_normalization_15[0][0]
-----
conv2d_16 (Conv2D) (None, 19, 19, 1024) 4718592
leaky_re_lu_15[0][0]
-----
batch_normalization_16 (BatchNo (None, 19, 19, 1024) 4096 conv2d_16[0][0]
-----
leaky_re_lu_16 (LeakyReLU) (None, 19, 19, 1024) 0
batch_normalization_16[0][0]
-----
conv2d_17 (Conv2D) (None, 19, 19, 512) 524288
leaky_re_lu_16[0][0]
-----
batch_normalization_17 (BatchNo (None, 19, 19, 512) 2048 conv2d_17[0][0]
-----
leaky_re_lu_17 (LeakyReLU) (None, 19, 19, 512) 0
batch_normalization_17[0][0]
-----
conv2d_18 (Conv2D) (None, 19, 19, 1024) 4718592
leaky_re_lu_17[0][0]
-----
batch_normalization_18 (BatchNo (None, 19, 19, 1024) 4096 conv2d_18[0][0]
-----
leaky_re_lu_18 (LeakyReLU) (None, 19, 19, 1024) 0
batch_normalization_18[0][0]
-----
conv2d_19 (Conv2D) (None, 19, 19, 1024) 9437184
leaky_re_lu_18[0][0]
-----
batch_normalization_19 (BatchNo (None, 19, 19, 1024) 4096 conv2d_19[0][0]

```



```

-----
conv2d_21 (Conv2D) (None, 38, 38, 64) 32768
leaky_re_lu_13[0][0]
-----
leaky_re_lu_19 (LeakyReLU) (None, 19, 19, 1024) 0
batch_normalization_19[0][0]
-----
batch_normalization_21 (BatchNo (None, 38, 38, 64) 256 conv2d_21[0][0]
-----
conv2d_20 (Conv2D) (None, 19, 19, 1024) 9437184
leaky_re_lu_19[0][0]
-----
leaky_re_lu_21 (LeakyReLU) (None, 38, 38, 64) 0
batch_normalization_21[0][0]
-----
batch_normalization_20 (BatchNo (None, 19, 19, 1024) 4096 conv2d_20[0][0]
-----
space_to_depth_x2 (Lambda) (None, 19, 19, 256) 0
leaky_re_lu_21[0][0]
-----
leaky_re_lu_20 (LeakyReLU) (None, 19, 19, 1024) 0
batch_normalization_20[0][0]
-----
concatenate_1 (Concatenate) (None, 19, 19, 1280) 0
space_to_depth_x2[0][0]
leaky_re_lu_20[0][0]
-----
conv2d_22 (Conv2D) (None, 19, 19, 1024) 11796480
concatenate_1[0][0]
-----
batch_normalization_22 (BatchNo (None, 19, 19, 1024) 4096 conv2d_22[0][0]
-----
leaky_re_lu_22 (LeakyReLU) (None, 19, 19, 1024) 0
batch_normalization_22[0][0]
-----

```

```
conv2d_23 (Conv2D)                (None, 19, 19, 425) 435625
leaky_re_lu_22[0][0]
```

```
=====
Total params: 50,983,561
Trainable params: 50,962,889
Non-trainable params: 20,672
-----
```

Feature extraction and conversion to bounding box params

```
[11]: # Retrieve features from the final layer of the model
      features = model.output

[12]: # Convert model features to bounding box parameters

      def features_to_boxes(features, anchors, num_classes):

          num_anchors = len(anchors)

          # Reshape to batch, height, width, num_anchors, box_params
          anchors_tensor = K.reshape(K.variable(anchors), [1, 1, 1, num_anchors, 2])

          # Dynamic implementation of conv dims for fully convolutional model
          conv_dims = K.shape(features)[1:3] # assuming channels last
          # In YOLO the height index is the inner most iteration
          conv_height_index = K.arange(0, stop=conv_dims[0])
          conv_width_index = K.arange(0, stop=conv_dims[1])
          conv_height_index = K.tile(conv_height_index, [conv_dims[1]])

          conv_width_index = K.tile(
              K.expand_dims(conv_width_index, 0), [conv_dims[0], 1])
          conv_width_index = K.flatten(K.transpose(conv_width_index))
          conv_index = K.transpose(K.stack([conv_height_index, conv_width_index]))
          conv_index = K.reshape(conv_index, [1, conv_dims[0], conv_dims[1], 1, 2])
          conv_index = K.cast(conv_index, K.dtype(features))

          features = K.reshape(
              features, [-1, conv_dims[0], conv_dims[1], num_anchors, num_classes + 5])
          conv_dims = K.cast(K.reshape(conv_dims, [1, 1, 1, 1, 2]), K.dtype(features))

          box_xy = K.sigmoid(features[..., :2])
          box_wh = K.exp(features[..., 2:4])
          box_confidence = K.sigmoid(features[..., 4:5])
          box_class_probs = K.softmax(features[..., 5:])
```

```

# Adjust predictions to each spatial grid point and anchor size.
# Note: YOLO iterates over height index before width index.
box_xy = (box_xy + conv_index) / conv_dims
box_wh = box_wh * anchors_tensor / conv_dims

```

```

return box_xy, box_wh, box_confidence, box_class_probs

```

```

[13]: yolo_outputs = features_to_boxes(features, anchors, len(class_names))

```

Evaluation of detected objects

Selection of most probable bounding boxes

```

[14]: # Evaluate detected objects

```

```

def yolo_evaluate(yolo_outputs, image_shape=(720., 1280.), max_output_size=10,
    score_threshold=.6, iou_threshold=.5):

    # Unpack bounding box params
    xy, wh, confidence, class_probs = yolo_outputs

    # Convert boxes information from (middle-point-coordinates, width-height) to
    corners coordinates
    boxes = boxes_to_corners(xy, wh)

    # Step 1: eliminate boxes with low probabilities - filter boxes with scores
    lower than the threshold
    scores, boxes, classes = filter_boxes(confidence, boxes, class_probs,
    score_threshold)

    # Step 2: eliminate boxes with non-max scores and high overlap - filter
    boxes with IoU lower than the threshold
    scores, boxes, classes = non_max_suppression(scores, boxes, classes,
    max_output_size, iou_threshold)

    # Scale boxes to match input image size
    boxes = scale_boxes(boxes, image_shape)

    return scores, boxes, classes

```

```

[15]: # Helper function: convert boxes to corners

```

```

def boxes_to_corners(box_xy, box_wh):

    box_mins = box_xy - (box_wh / 2.)
    box_maxes = box_xy + (box_wh / 2.)

```

```

return K.concatenate([
    box_mins[..., 1:2], # y_min
    box_mins[..., 0:1], # x_min
    box_maxes[..., 1:2], # y_max
    box_maxes[..., 0:1]]) # x_max

```

[16]: *# Helper function: eliminate boxes data with low predicted probability*

```

def filter_boxes(confidence, boxes, class_probs, threshold = 0.6):

    # Compute the score for each box (object probability multiplied by class
    →probability)
    scores = confidence * class_probs

    # Find class with maximum score for each box and respective score
    classes_with_max_score = K.argmax(scores, axis=-1)
    max_scores = K.max(scores, axis=-1)

    # Define a mask (binary filter) to discard boxes with score lower than the
    →set threshold
    mask = (max_scores >= threshold)

    # Apply the mask to discard boxes data with the score lower than the set
    →threshold
    scores = tf.boolean_mask(max_scores, mask)
    boxes = tf.boolean_mask(boxes, mask)
    classes = tf.boolean_mask(classes_with_max_score, mask)

    return scores, boxes, classes

```

[17]: *'''*
Calculate intersection over union ratio. This function is not used in this
→implementation.
Here just as an exercise.
'''

```

def iou(box1, box2):

    # Unpack vertices information
    (x1_1, y1_1, x2_1, y2_1) = box1
    (x1_2, y1_2, x2_2, y2_2) = box2

    # Find area of intersection
    width = min(x2_1, x2_2) - max(x1_1, x1_2)
    height = min(y2_1, y2_2) - max(y1_1, y1_2)
    intersection = max(width, 0) * max(height, 0)

```

```

# Find union area
area_1 = (x2_1 - x1_1) * (y2_1 - y1_1)
area_2 = (x2_2 - x1_2) * (y2_2 - y1_2)
union = area_1 + area_2 - intersection

```

```

# Compute IoU
iou = intersection / union

```

```

return iou

```

[18]: # Helper function: perform non-max suppression to make sure algorithm detects
↳ objects only once

```

def non_max_suppression(scores, boxes, classes, max_output_size = 10,
↳ iou_threshold = 0.5):

    # Initialize max_output_size variable (we need it later to perform tf.image.
↳ non_max_suppression)
    max_output_size_tensor = K.variable(max_output_size, dtype='int32')
    K.get_session().run(tf.variables_initializer([max_output_size_tensor]))

    # Return indices from the boxes data to keep: eliminate non-max scores with
↳ a significant overlap with the max score
    remaining_indices = tf.image.non_max_suppression(boxes, scores,
↳ max_output_size_tensor, iou_threshold=0.5)

    # Get rid of bounding box data with non-max scores and significant overlap
↳ with max score
    scores = K.gather(scores, remaining_indices)
    boxes = K.gather(boxes, remaining_indices)
    classes = K.gather(classes, remaining_indices)

    return scores, boxes, classes

```

[19]: # Helper function: Scale predicted boxes to make them drawable on the image
↳ (normalize boxes by image dimensions)

```

def scale_boxes(boxes, image_shape):

    # Retrieve image dimensions
    height = image_shape[0]
    width = image_shape[1]

    # Stack and reshape image dimensions to correspond to boxes coordinates
↳ (b_x, b_y, b_h, b_w)

```

```

image_dims = K.stack([height, width, height, width])
image_dims = K.reshape(image_dims, [1, 4])
image_dims = K.cast(image_dims, 'float32')
# Normalize boxes by image dimensions
boxes = boxes * image_dims

return boxes

```

Testing

[20]: *# Make predictions and draw predicted bounding boxes on the input image*

```

def predict(sess, img_path):

    # Preprocess input images to match model input size
    image, image_data, image_shape = preprocess_image(img_path, model_image_size,
    ↪= (608, 608))

    # Optimize: perform forward propagation
    out_scores, out_boxes, out_classes = sess.run(yolo_evaluate(yolo_outputs,
    ↪image_shape), feed_dict={model.input: image_data,

    ↪                                K.learning_phase(): 0})

    # Draw bounding boxes on the input image
    draw_boxes(image, out_scores, out_boxes, out_classes, class_names)

    # Save edited image
    path, filename = os.path.split(img_path)
    image.save(os.path.join('./out', filename), quality=90)

    # Display the image
    output_image = imageio.imread(os.path.join('out/', filename))
    imshow(output_image)

    return out_scores, out_boxes, out_classes

```

[21]: *# Helper function: draw selected bounding boxes on the image*

```

def draw_boxes(image, out_scores, out_boxes, out_classes, class_names):

    # Generate colors for drawing predicted bounding boxes
    colors = generate_colors(class_names)

    # Load a font - workaround for ImageFont.truetype OSError
    file = open("./font/FiraMono-Medium.otf", "rb")
    bytes_font = BytesIO(file.read())

```

```

font = ImageFont.truetype(bytes_font, size=np.floor(3e-2 * image.size[1] + 0.
→5).astype('int32'))
file.close()

# Set bounding box line thickness
thickness = (image.size[0] + image.size[1]) // 300

for i, c in reversed(list(enumerate(out_classes))):

    # Retrieve box parameters
    predicted_class = class_names[c]
    box = out_boxes[i]
    score = out_scores[i]

    # Create output label
    label = '{: {:.2f}}'.format(predicted_class, score)

    # Draw the bounding box on the image
    draw = ImageDraw.Draw(image)
    label_size = draw.textsize(label, font)

    # Retrieve box coordinates
    top, left, bottom, right = box

    # Limit bounding box coordinates to the image dimensions
    top = max(0, np.floor(top + 0.5).astype('int32'))
    left = max(0, np.floor(left + 0.5).astype('int32'))
    bottom = min(image.size[1], np.floor(bottom + 0.5).astype('int32'))
    right = min(image.size[0], np.floor(right + 0.5).astype('int32'))


    # Set origin for the label text
    if top - label_size[1] >= 0:
        text_origin = np.array([left, top - label_size[1]])
    else:
        text_origin = np.array([left, top + 1])

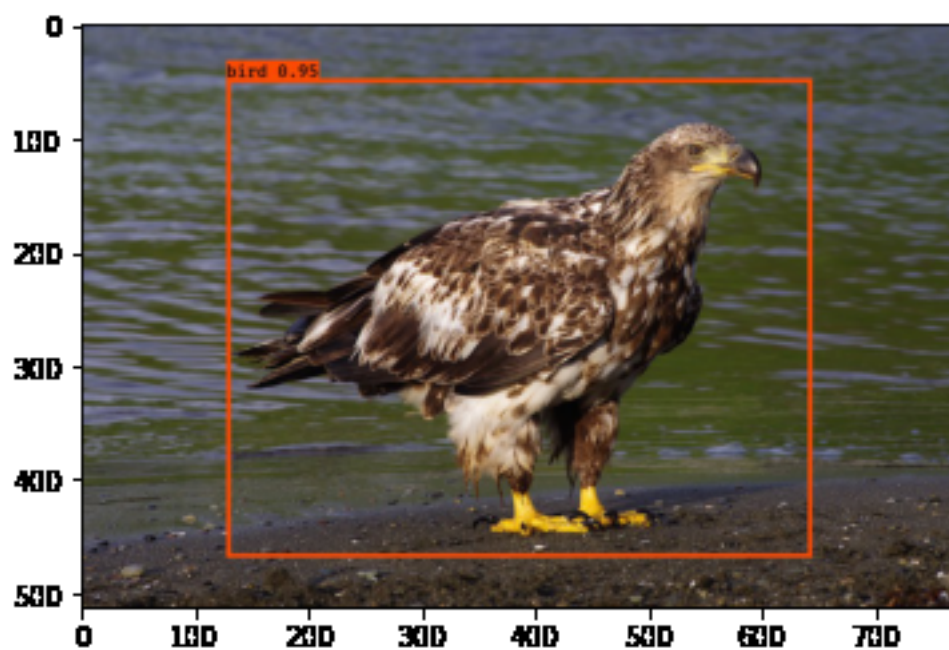
    # My kingdom for a good redistributable image drawing library.
    for i in range(thickness):
        draw.rectangle([left + i, top + i, right - i, bottom - i],
→outline=colors[c])
        draw.rectangle([tuple(text_origin), tuple(text_origin + label_size)],
→fill=colors[c])
        draw.text(text_origin, label, fill=(0, 0, 0), font=font)
    del draw

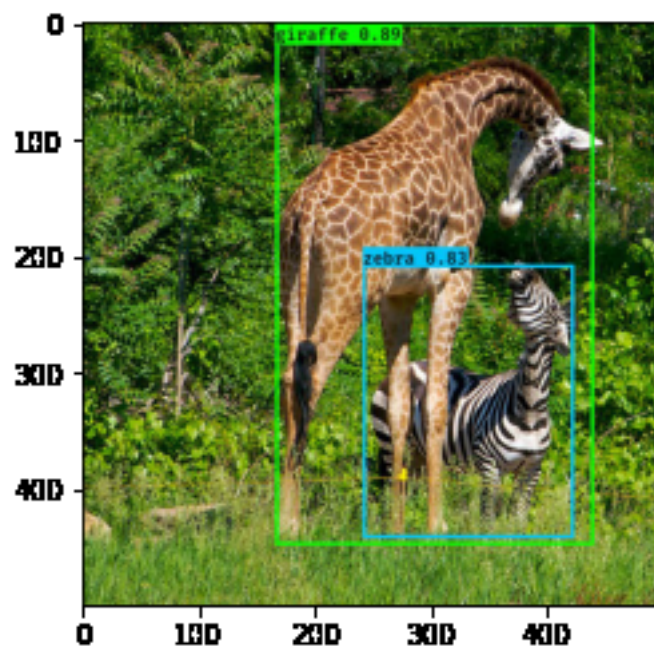

```

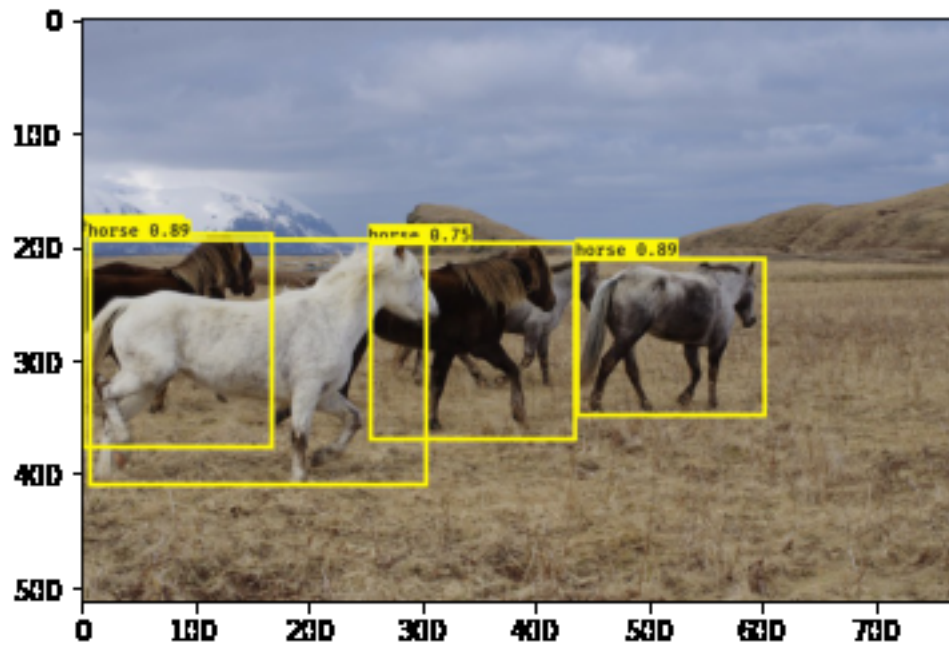
```
[22]: for filename in os.listdir('./yad2k/images'):
      if filename.endswith('.jpg'):
          predict(sess, './yad2k/images/' + str(filename))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/site-packages/tensorflow/python/ops/array_ops.py:1354: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where













```
[23]: for filename in os.listdir('./images'):
      if filename.endswith('.jpg'):
          predict(sess, './images/' + str(filename))
```







[]: